

d-comparing-the-different-ml-model

October 20, 2024

1 Website traffic prediction and comparing the different ML model

1.1 Rahul Manjhi

```
[73]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression
```

2 Dataset Exploration:

```
[ ]:
```

```
[13]: data = pd.read_csv("website_traffic.csv")
```

```
[15]: data
```

```
[15]:
```

	Page Views	Session Duration	Bounce Rate	Traffic Source	Time on Page \
0	5	11.051381	0.230652	Organic	3.890460
1	4	3.429316	0.391001	Social	8.478174
2	4	1.621052	0.397986	Organic	9.636170
3	5	3.629279	0.180458	Organic	2.071925
4	5	4.235843	0.291541	Paid	1.960654
...
1995	1	2.724513	0.207187	Referral	1.324206
1996	3	0.392856	0.095559	Organic	3.824416
1997	4	9.899823	0.446622	Organic	1.288675
1998	3	0.393319	0.278340	Paid	5.037584
1999	3	0.882638	0.338026	Direct	5.186908

	Previous Visits	Conversion Rate
0	3	1.0
1	0	1.0
2	2	1.0
3	3	1.0
4	5	1.0
...
1995	2	1.0
1996	1	1.0
1997	1	1.0
1998	2	1.0
1999	3	1.0

[2000 rows x 7 columns]

```
[17]: missing = data.isnull()
```

```
[19]: missing
```

```
[19]:
```

	Page Views	Session Duration	Bounce Rate	Traffic Source	Time on Page \
0	False	False	False	False	False
1	False	False	False	False	False
2	False	False	False	False	False
3	False	False	False	False	False
4	False	False	False	False	False
...
1995	False	False	False	False	False
1996	False	False	False	False	False
1997	False	False	False	False	False
1998	False	False	False	False	False
1999	False	False	False	False	False

	Previous Visits	Conversion Rate
0	False	False
1	False	False
2	False	False
3	False	False
4	False	False
...
1995	False	False
1996	False	False
1997	False	False
1998	False	False
1999	False	False

[2000 rows x 7 columns]

```
[21]: coun_missing_values = data.isnull().sum()
```

```
[23]: coun_missing_values
```

```
[23]: Page Views      0
      Session Duration 0
      Bounce Rate     0
      Traffic Source   0
      Time on Page     0
      Previous Visits  0
      Conversion Rate  0
      dtype: int64
```

```
[25]: summary_statistics = data.describe()
```

```
[27]: summary_statistics
```

```
[27]:
```

	Page Views	Session Duration	Bounce Rate	Time on Page	\
count	2000.000000	2000.000000	2000.000000	2000.000000	
mean	4.950500	3.022045	0.284767	4.027439	
std	2.183903	3.104518	0.159781	2.887422	
min	0.000000	0.003613	0.007868	0.068515	
25%	3.000000	0.815828	0.161986	1.935037	
50%	5.000000	1.993983	0.266375	3.315316	
75%	6.000000	4.197569	0.388551	5.414627	
max	14.000000	20.290516	0.844939	24.796182	

	Previous Visits	Conversion Rate
count	2000.000000	2000.000000
mean	1.978500	0.982065
std	1.432852	0.065680
min	0.000000	0.343665
25%	1.000000	1.000000
50%	2.000000	1.000000
75%	3.000000	1.000000
max	9.000000	1.000000

```
[29]: coun_missing_values, summary_statistics
```

```
[29]: (Page Views      0
      Session Duration 0
      Bounce Rate     0
      Traffic Source   0
      Time on Page     0
      Previous Visits  0
      Conversion Rate  0
      dtype: int64,
```

	Page Views	Session Duration	Bounce Rate	Time on Page	\
count	2000.000000	2000.000000	2000.000000	2000.000000	
mean	4.950500	3.022045	0.284767	4.027439	
std	2.183903	3.104518	0.159781	2.887422	
min	0.000000	0.003613	0.007868	0.068515	
25%	3.000000	0.815828	0.161986	1.935037	
50%	5.000000	1.993983	0.266375	3.315316	
75%	6.000000	4.197569	0.388551	5.414627	
max	14.000000	20.290516	0.844939	24.796182	

	Previous Visits	Conversion Rate
count	2000.000000	2000.000000
mean	1.978500	0.982065
std	1.432852	0.065680
min	0.000000	0.343665
25%	1.000000	1.000000
50%	2.000000	1.000000
75%	3.000000	1.000000
max	9.000000	1.000000

```
[ ]:
```

```
[32]: # Perform one-hot encoding on the 'Traffic Source' categorical variable
df_encoded = pd.get_dummies(data, columns=['Traffic Source'], drop_first=True)
```

```
[34]: df_encoded.head(10)
```

```
[34]:
```

	Page Views	Session Duration	Bounce Rate	Time on Page	Previous Visits	\
0	5	11.051381	0.230652	3.890460	3	
1	4	3.429316	0.391001	8.478174	0	
2	4	1.621052	0.397986	9.636170	2	
3	5	3.629279	0.180458	2.071925	3	
4	5	4.235843	0.291541	1.960654	5	
5	3	4.541868	0.420740	3.438712	2	
6	5	1.949558	0.034978	2.119271	1	
7	4	1.685740	0.252343	3.478016	5	
8	6	0.033268	0.120703	5.285519	1	
9	7	7.833742	0.212727	4.060115	5	

	Conversion Rate	Traffic Source_Organic	Traffic Source_Paid	\
0	1.0	True	False	
1	1.0	False	False	
2	1.0	True	False	
3	1.0	True	False	
4	1.0	False	True	
5	1.0	False	False	
6	1.0	False	False	

7	1.0	False	True
8	1.0	True	False
9	1.0	False	True

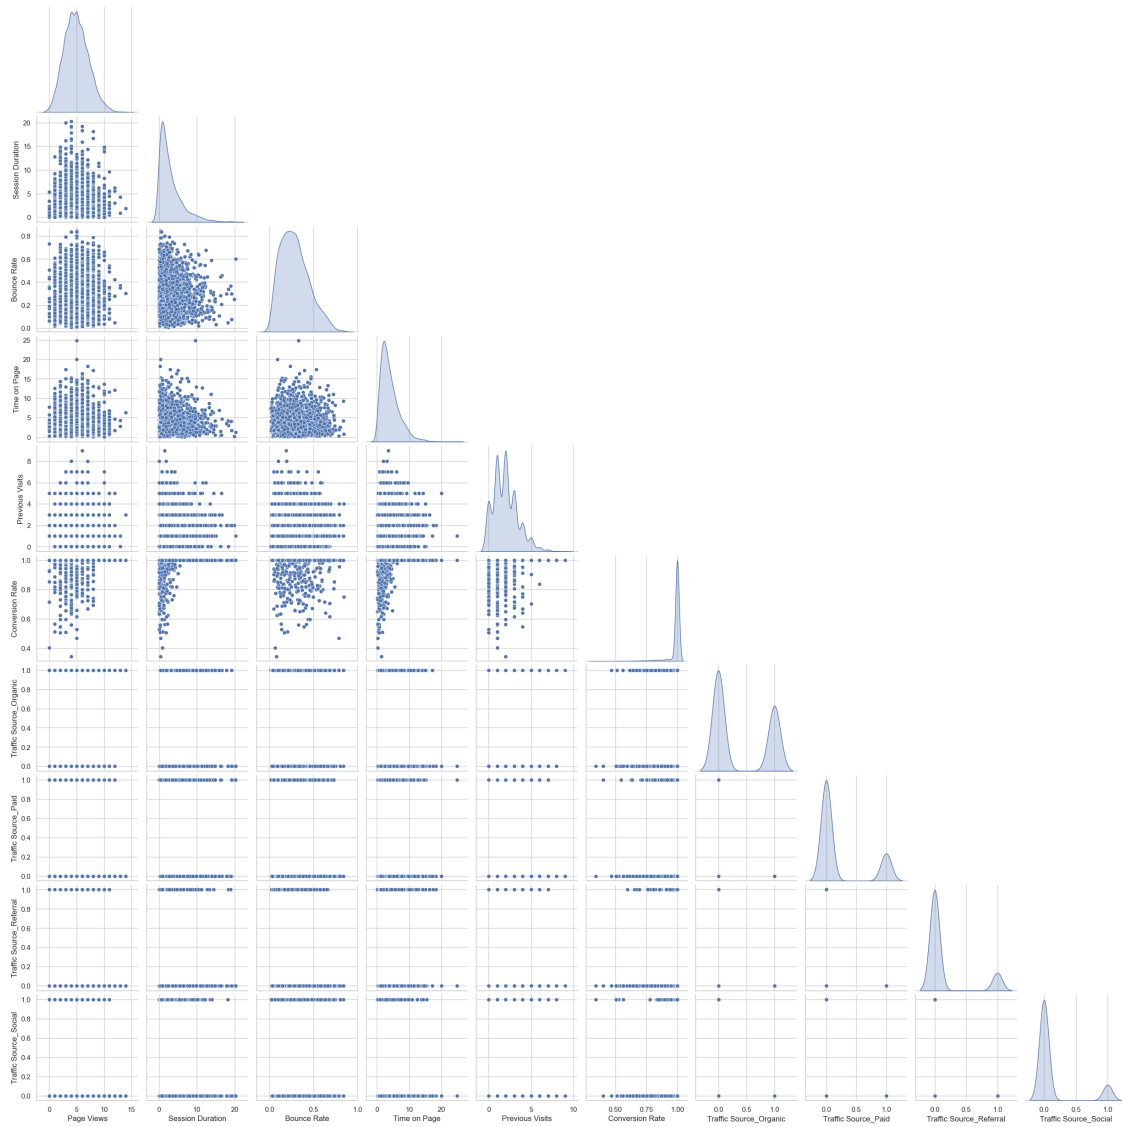
	Traffic Source_Referral	Traffic Source_Social
0	False	False
1	False	True
2	False	False
3	False	False
4	False	False
5	False	True
6	False	True
7	False	False
8	False	False
9	False	False

```
[ ]:
```

```
[37]: # Set up the plotting environment
sns.set(style="whitegrid")
```

```
[39]: # Scatter plot matrix (pair plot) to visualize relationships between numeric_
      ↪ variables
sns.pairplot(df_encoded, diag_kind='kde', corner=True)
plt.suptitle('Pair Plot of Website Traffic Data', y=1.02)
plt.show()
```

Pair Plot of Website Traffic Data



[]:

[]:

3 Multiple Linear Regression:

```
[56]: # Define the features (X) and target (y)
X = df_encoded.drop(columns=['Conversion Rate'])
y = df_encoded['Conversion Rate']
```

```
[58]: # Split the data into 80% training and 20% testing
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

```
# Instantiate the Multiple Linear Regression model
model = LinearRegression()
```

```
[60]: # Train the model on the training data
model.fit(X_train, y_train)
```

```
[60]: LinearRegression()
```

```
[62]: y_pred = model.predict(X_test)
```

```
[ ]:
```

```
[65]: # Evaluate the model
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)
```

```
[67]: # Calculate adjusted R-squared
n = X_test.shape[0]
p = X_test.shape[1]
adjusted_r2 = 1 - (1 - r2) * (n - 1) / (n - p - 1)

mae, mse, rmse, r2, adjusted_r2
```

```
[67]: (0.03163385542916112,
0.0027909306686198193,
0.0528292595880334,
0.10626266875894574,
0.0856379611149215)
```

```
[ ]:
```

```
[ ]:
```

4 Polynomial Regression

```
[75]: # Polynomial transformation (degree 2 or 3)
degree = 2 # or 3 for higher-degree polynomial
poly = PolynomialFeatures(degree=degree)
X_poly_train = poly.fit_transform(X_train)
X_poly_test = poly.transform(X_test)
```

[]:

```
[77]: # Fit the linear regression model with polynomial features
model_poly = LinearRegression()
model_poly.fit(X_poly_train, y_train)

# Predict on the test set
y_pred_poly = model_poly.predict(X_poly_test)
```

[]:

```
[79]: # Evaluate the model
mae_poly = mean_absolute_error(y_test, y_pred_poly)
mse_poly = mean_squared_error(y_test, y_pred_poly)
rmse_poly = np.sqrt(mse_poly)
r2_poly = r2_score(y_test, y_pred_poly)
```

[]:

```
[81]: # Calculate adjusted R-squared
n = X_test.shape[0]
p = X_poly_test.shape[1] # Account for the new number of features in the
    ↪ polynomial model
adjusted_r2_poly = 1 - (1 - r2_poly) * (n - 1) / (n - p - 1)
```

[]:

```
[83]: # Output results
print("Polynomial Regression (degree {}):".format(degree))
print(f"MAE: {mae_poly}")
print(f"MSE: {mse_poly}")
print(f"RMSE: {rmse_poly}")
print(f"R²: {r2_poly}")
print(f"Adjusted R²: {adjusted_r2_poly}")
```

```
Polynomial Regression (degree 2):
MAE: 0.02912565018405664
MSE: 0.0024035453442120837
RMSE: 0.04902596602018244
R²: 0.2303147385186486
Adjusted R²: 0.10725459496785117
```

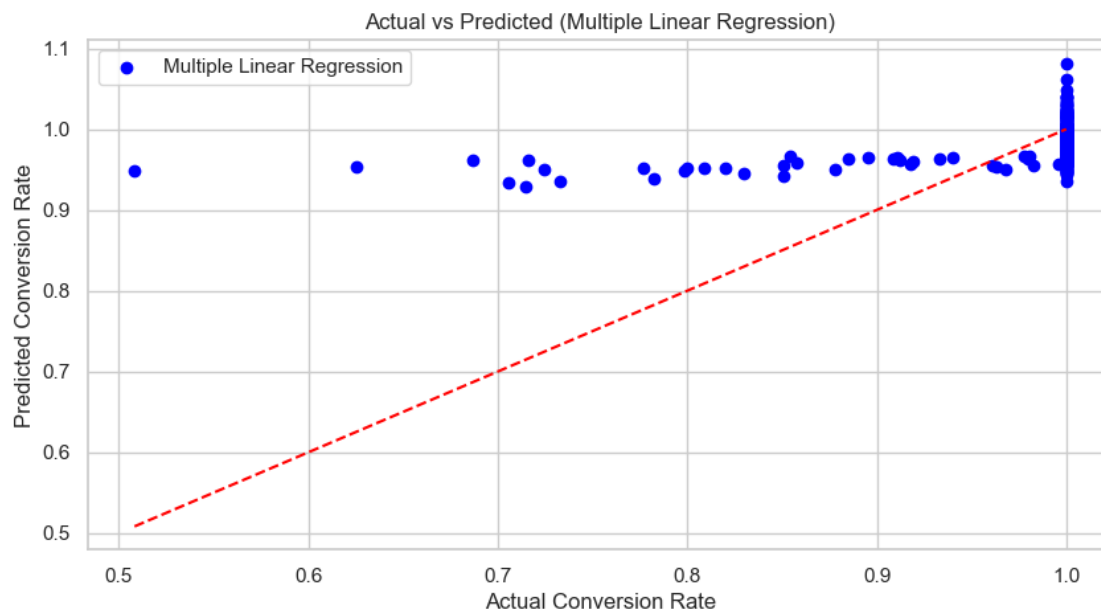
[]:

[]:

5 Insights and Reporting

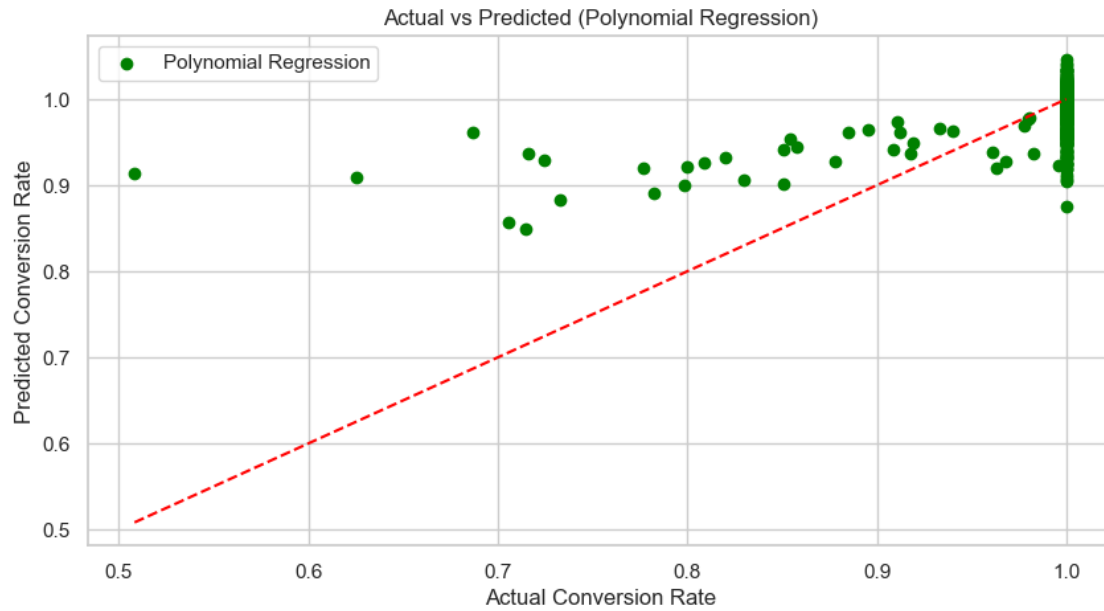
5.1 1. Plotting Actual vs Predicted Values

```
[86]: # Plot actual vs predicted for Multiple Linear Regression
plt.figure(figsize=(10,5))
plt.scatter(y_test, y_pred, color='blue', label='Multiple Linear Regression')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red',
         linestyle='--')
plt.xlabel('Actual Conversion Rate')
plt.ylabel('Predicted Conversion Rate')
plt.title('Actual vs Predicted (Multiple Linear Regression)')
plt.legend()
plt.show()
```



```
[ ]:
```

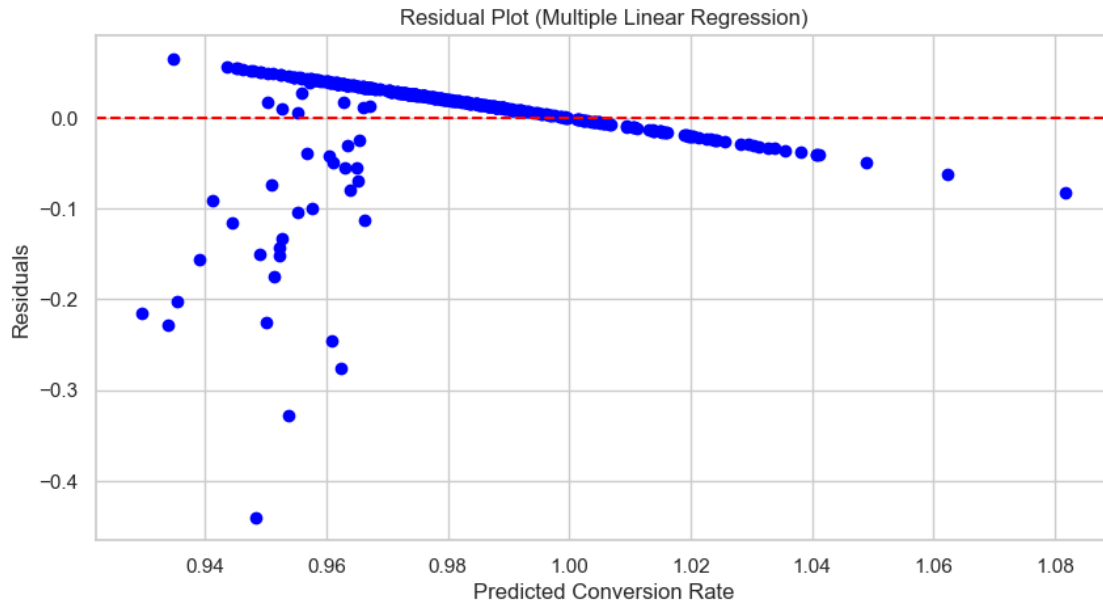
```
[88]: # Plot actual vs predicted for Polynomial Regression
plt.figure(figsize=(10,5))
plt.scatter(y_test, y_pred_poly, color='green', label='Polynomial Regression')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], color='red',
         linestyle='--')
plt.xlabel('Actual Conversion Rate')
plt.ylabel('Predicted Conversion Rate')
plt.title('Actual vs Predicted (Polynomial Regression)')
plt.legend()
plt.show()
```



[]:

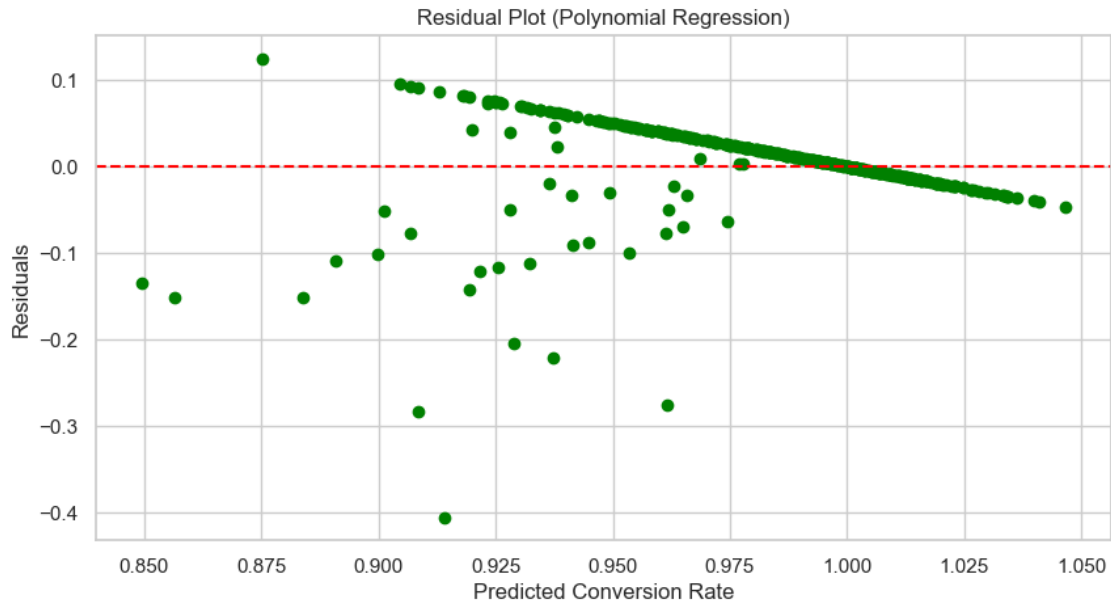
5.2 2. Residual Plot

```
[92]: # Residual plot for Multiple Linear Regression
residuals = y_test - y_pred
plt.figure(figsize=(10,5))
plt.scatter(y_pred, residuals, color='blue')
plt.axhline(y=0, color='red', linestyle='--')
plt.xlabel('Predicted Conversion Rate')
plt.ylabel('Residuals')
plt.title('Residual Plot (Multiple Linear Regression)')
plt.show()
```



[]:

```
[94]: # Residual plot for Polynomial Regression
residuals_poly = y_test - y_pred_poly
plt.figure(figsize=(10,5))
plt.scatter(y_pred_poly, residuals_poly, color='green')
plt.axhline(y=0, color='red', linestyle='--')
plt.xlabel('Predicted Conversion Rate')
plt.ylabel('Residuals')
plt.title('Residual Plot (Polynomial Regression)')
plt.show()
```

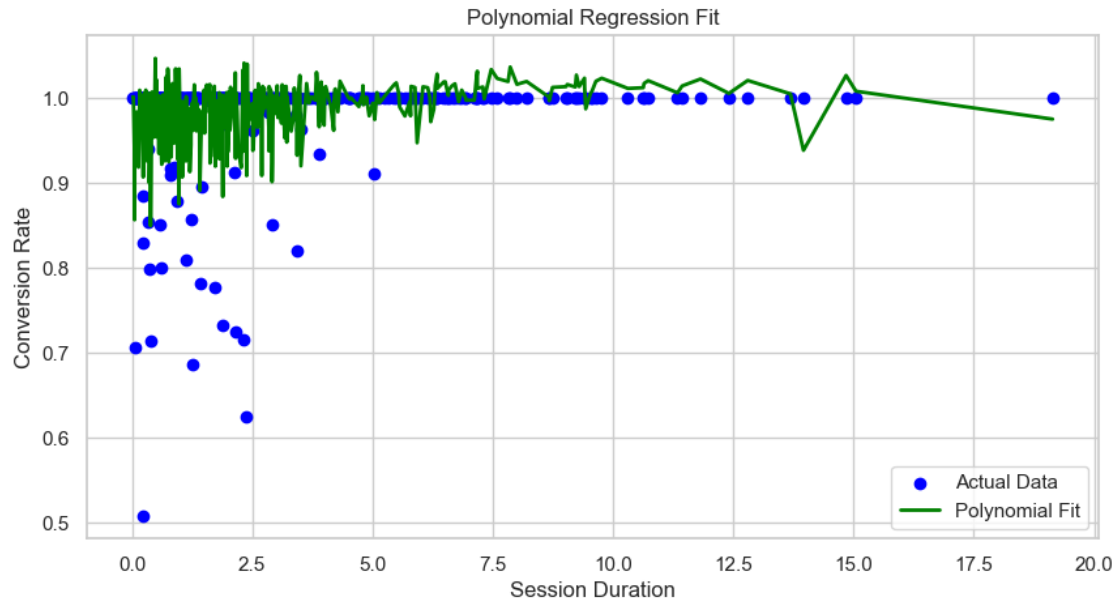


[]:

5.3 3. Plot Polynomial Curve

```
[97]: # Sort the test data to get a smooth curve
sorted_indices = np.argsort(X_test['Session Duration'])
X_test_sorted = X_test['Session Duration'].values[sorted_indices]
y_test_sorted = y_test.values[sorted_indices]
y_pred_poly_sorted = y_pred_poly[sorted_indices]

[99]: # Plot polynomial regression curve
plt.figure(figsize=(10,5))
plt.scatter(X_test_sorted, y_test_sorted, color='blue', label='Actual Data')
plt.plot(X_test_sorted, y_pred_poly_sorted, color='green', label='Polynomial_
↳Fit', linewidth=2)
plt.xlabel('Session Duration')
plt.ylabel('Conversion Rate')
plt.title('Polynomial Regression Fit')
plt.legend()
plt.show()
```



[]:

6 Conclusion and Reporting:

6.0.1 1 Summary of Findings:

Compare the metrics (MAE, MSE, RMSE, R^2 , and Adjusted R^2) between the multiple linear regression and polynomial regression models. If polynomial regression has significantly lower errors and a better fit, it suggests that the relationships in the data are non-linear, and polynomial features capture these better.

[]:

2 Next Steps for Analysis and Model Building: Feature Engineering: You could explore higher-degree polynomial features or interaction terms.

Regularization: Consider using Ridge or Lasso regression to handle potential overfitting in polynomial models.

Cross-Validation: To ensure model stability, use k-fold cross-validation for better evaluation of model performance.

Other Algorithms: Try more complex models like decision trees, random forests, or gradient boosting if polynomial regression performs significantly better than linear models.